

# Probabilistic Inference by Hashing and Optimization

Focusing on Approximate Model Counting

Stefano Ermon  
slides by Dor Cohen

June 13, 2017

# Introduction

- ▶ Many problems in ML and Statistics involve computation of high-dimensional integrals (e.g. evaluate posterior probabilities)
- ▶ Computing expectations *w.r.t* high dimensional probability distributions known to be intractable in worst-case (Roth, 2016)
- ▶ Difficulty arises as number of possible states scales exponentially - *curse of dimensionality* (Bellman, 1996)

# Standard approaches

Focusing on *discrete* probability distributions and *approximately* computing expectations, there are two main standard approaches:

- ▶ **Monte Carlo** sampling techniques (1950s): Approximate complex distributions with small number of representative samples
- ▶ **Variational** methods: Approximate complex models using families of tractable distributions (e.g., estimate  $P(Z|X) \approx Q(Z)$  and compare distributions using *KL – divergence*)

These don't provide tight guarantees on accuracy

# MCMC - Monte Carlo Markov Chain

**Goal:** Sample from a *discrete* distributions vector  $\pi = (\pi_1, \dots, \pi_n)$

**Idea:**

- ▶ Build a Markov chain with  $n$  states such that it will be *ergodic*
- ▶ Choose weights for graph edges (probabilities) such that the chain will be *invertible w.r.t* to the desirable distribution  $\pi$
- ▶ *ergodic* and *invertible* chain imply that the "last" state of long random walk on the chain is distributed like  $\pi$

# MCMC - Monte Carlo Markov Chain

## Key difficulty

is to to draw proper samples

- ▶ Needs to set up a Markov Chain over entire state space
- ▶ Has to reach equilibrium distribution - requires exponential time (Madras, 2002)
- ▶ In practice will only give approximate answer. Chain may "trap" in less relevant areas of the state space

## Convergence of *Metropolis – Hastings*

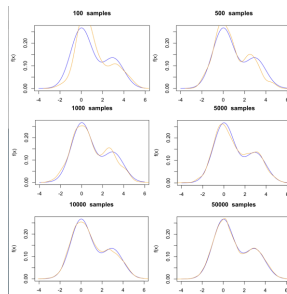


Figure: MCMC attempts to approximate blue distribution with orange one [Wiki]

# A new approach


Approximate inference/counting algorithms based on randomized hashing and optimizations:

- ▶ *provably accurate results* assuming access to an **optimization oracle**
- ▶ Can compute parity functions, marginal probabilities providing an approximately correct answer
- ▶ Specifically, within a factor of  $1 + \epsilon$  of the true value for any desired  $\epsilon$ , with high probability (at least  $1 - \lambda$  for any desired  $\lambda > 0$ )

# Hashing and optimization

Statistics are still computed using small subset of samples.  
 However these samples aren't drawn randomly from the distribution using an MC, but:

- ▶ Samples are obtained by random projection of original high-dimensional states to lower-dimensional ones using *Universal hash functions*
- ▶ Then look for "Most likely" configuration in the projected subspace (current optimization tools can handle millions of variables)




# Complexity perspective


Inference or counting problems we consider are in  $\#P$ , believed to be significantly harder than  $NP$  (Valiant, 1979)

- ▶ **If we allow small failure probability, and small error** we can *reduce* the  $\#P$  problem to smaller number of instances of a *NP-equivalent* (combinatorial) optimization problem

**Integration:** How many ways to set the variables s.t. a certain property holds?



**Search:** Is there a way to set the variables such that some property holds?





## Problem statement - setup

- ▶ Let  $\Sigma$  be a large but finite set (e.g. the set of all possible assignments in a model) and let  $w : \Sigma \rightarrow \mathbb{R}^+$  be a non-negative function that assigns a weight for each configuration (an element of  $\Sigma$ )
- ▶ We are given  $2^n$  configurations, and non-negative weights  $w$  (Bayes net, Factor graph, weighted CNF)
- ▶ **Goal:** (approximately) compute the total weight of the set, defined as:

$$W = \sum_{\sigma \in \Sigma} w(\sigma)$$

- ▶ **Example:**  $n = 100$  variables, sum over  $2^{100}$  configurations

# Problem statement - setup

**Assumption:** We assume that we have access to an *optimization oracle* that can solve the following optimization problem:

$$\max_{\sigma \in \Sigma} w(\sigma) \mathbb{1}_{\{o\}}(\sigma)$$

Where  $\mathbb{1}_{\{o\}} : \Sigma \rightarrow \{0, 1\}$  is an indicator function for compactly represented subset  $o \subseteq \Sigma$ , i.e.,  $\mathbb{1}_{\{o\}} = 1$  if  $\sigma \in o$  and 0 otherwise.

- ▶  $o$  may be compactly represented using a smaller subset of constraints

# SAT

- ▶ A boolean formula  $\varphi$  is said to be in *CNF* if its a logical conjunction of a set of clauses  $C_1, \dots, C_n$ , where each clause  $C$  is a logical disjunction of a set of literals. e.g.,  $(x_1 \vee \neg x_2)$
- ▶ *SAT*: deciding if there exists an assignment that **satisfies**  $\varphi$

**Example:**  $\varphi = (x_1 \vee \neg x_2 \vee x_3) \wedge (x_4 \vee \neg x_1) \wedge (x_2 \vee \neg x_3)$

*Satisfying assignment:*  $\{x_1 = 1, x_2 = 1, x_3 = 1, x_4 = 1\}$

$\implies \varphi$  is SATISFIABLE

# Model counting

- ▶ Let  $V$  be the set of boolean variables of  $\varphi$ ,  $|V| = n$ , and let  $\sigma = \{0, 1\}^n$  be the set of all possible assignments to these variables
- ▶ An assignment  $\sigma \in \Sigma$  is a mapping that assigns a value in  $\{0, 1\}$  to each variable in  $V$
- ▶ Define the weight  $w(\sigma)$  to be 1 if  $\varphi$  is satisfied by  $\sigma$  and 0 otherwise

In this context:

$$W = \sum_{\sigma \in \Sigma} w(\sigma) = \#(\varphi)$$

# Approximate model counting

- ▶ Problem of *approximately* counting the number of solutions, assuming access to an *NP* oracle (e.g. SAT solver) was first considered by Stockmeyer (1985) - Important result he established:  $\#P$  can be approximated in  $BPP^{NP}$ <sup>1</sup>

Intuition behind his algorithm:

- ▶ Let  $S \subseteq \Sigma$  be the set of solutions to  $\varphi$
- ▶ Reliably estimate  $|S|$  by **repeating** the following process
  1. Randomly partition  $\Sigma$  into  $2^m$  cells
  2. Select one of the cells
  3. Compute if  $S$  has at least one element in this cell (invoke a SAT solver)

<sup>1</sup>

<sup>1</sup>algorithms that have bounded-error probabilistic polynomial time and access to an *NP* oracle



# Approximate model counting

To summarize:

- ▶ Estimate  $|S|$ , by defining progressively smaller cells, until no element of  $S$  can be found inside a randomly chosen cell
- ▶ The larger  $|S|$  is, the smaller the cells have to be
- ▶ **Correctness** of the algorithm relies crucially on how the space is randomly partitioned into cells
- ▶ To achieve strong worst-case (probabilistic) guarantees, algorithm relies on *universal hash functions* (Vadhan, 2011; Goldreich, 2011)

# Universal hash functions

**Definition:** A family of functions  $\mathcal{H} = \{h : \{0, 1\}^n \mapsto \{0, 1\}^m\}$  is *strongly universal* (pairwise independent) if the following two conditions hold when  $h$  is chosen uniformly at random from  $\mathcal{H}$  :

1.  $\forall x \in \{0, 1\}^n$ , the random variable  $H(x)$  is uniformly distributed in  $\{0, 1\}^m$
2.  $\forall x_1, x_2 \in \{0, 1\}^n \mid x_1 \neq x_2$ , the random variables  $H(x_1)$  and  $H(x_2)$  are independent



# Universal hash functions

- ▶ Considering the set of all possible functions from  $\{0, 1\}^n$  to  $\{0, 1\}^m$  we establish statistically optimal functions. Easy to verify this is a family of *fully* independent functions
- ▶ However, functions like these require  $m2^n$  bits for specifying  
→ not useful for large  $n$
- ▶ *pairwise independent* hash functions can be specified compactly, with number of bits linear in  $n$

# Pairwise independent hash functions

- Generally, these functions are based on modular arithmetic constraints of the form  $Ax = b \bmod 2$ . Implies that  $Ax$  is congruent to  $b$  modulo 2

**Proposition:** Let  $A \in \{0, 1\}^{m \times n}$ ,  $b \in \{0, 1\}^m$ . The family  $\mathcal{H} = \{h_{A,b}(x) : \{0, 1\}^n \mapsto \{0, 1\}^m\}$  where  $h_{A,b}(x) = Ax + b \bmod 2$  is a family of pairwise independent hash functions.

# Algorithm: ApproxModelCount

---

**Algorithm 9.1** ApproxModelCount ( $F, \mathcal{O}_S, \Delta$ )

---

```
1: Let  $S$  denote the set of solutions to the input formula  $F$ 
2:  $T \leftarrow \left\lceil \frac{\log(1/\Delta)}{\alpha} \log n \right\rceil$ 
3:  $i = 0$ 
4: while  $i \leq n$  do
5:   for  $t = 1, \dots, T$  do
6:     Sample hash function  $h_{A,b}^i : \Sigma \rightarrow \{0, 1\}^i$ , i.e.
7:     sample uniformly  $A \in \{0, 1\}^{i \times n}$ ,  $b \in \{0, 1\}^i$ 
8:     Let  $S(h_{A,b}^i) = |\{x \in S \mid Ax \equiv b \pmod{2}\}|$ 
9:      $w_i^t \leftarrow \mathbb{I}[S(h_{A,b}^i) \geq 1]$ , using  $\mathcal{O}_S$  to check if  $\{x \in S \mid Ax \equiv b \pmod{2}\}$  is empty
10:   end for
11:   if Median( $w_i^1, \dots, w_i^T$ )  $< 1$  then
12:     break
13:   end if
14:    $i = i + 1$ 
15: end while
16: Return  $\lfloor 2^{i-1} \rfloor$ 
```

---

# Algorithm: ApproxModelCount

Based on pairwise independence, possible to show that ApproxModelCount provides with high probability an accurate estimate of the true value, summarized by the following Theorem:

**Theorem:** For any  $\Delta > 0$ , positive constant  $a \leq 0.0042$ , ApproxModelCount makes  $\Theta(n \ln n \ln 1/\delta)$  queries to the  $NP$  oracle  $O_S$  and, with probability of at least  $(1 - \Delta)$ , outputs a 16-approximation of  $|S|$ , the number of solutions of a formula  $\varphi$



# Theorem proof sketch

By linearity of expectation:  
 $E[S(h_{A,b}^i)] = \frac{|S|}{2^i}$ . It implies that *in expectation* the while loop should break for  $i \approx \log|S|$  i.e., when the corresponding expected number of "surviving" solutions is less than 1. When this happens, the algorithm provides an accurate estimate for  $|S|$

---

## Algorithm 9.1 ApproxModelCount ( $F, \mathcal{O}_S, \Delta$ )

---

```

1: Let  $S$  denote the set of solutions to the input formula  $F$ 
2:  $T \leftarrow \left\lceil \frac{\log(1/\Delta)}{\alpha} \log n \right\rceil$ 
3:  $i = 0$ 
4: while  $i \leq n$  do
5:   for  $t = 1, \dots, T$  do
6:     Sample hash function  $h_{A,b}^i : \Sigma \rightarrow \{0,1\}^i$ , i.e.
7:     sample uniformly  $A \in \{0,1\}^{i \times n}$ ,  $b \in \{0,1\}^i$ 
8:     Let  $S(h_{A,b}^i) = \{x \in S \mid Ax \equiv b \pmod{2}\}$ 
9:      $w_t^i \leftarrow \mathbb{I}[S(h_{A,b}^i) \geq 1]$ , using  $\mathcal{O}_S$  to check if  $\{x \in S \mid Ax \equiv b \pmod{2}\} \neq \emptyset$ 
10:   end for
11:   if  $\text{Median}(w_1^i, \dots, w_T^i) < 1$  then
12:     break
13:   end if
14:    $i = i + 1$ 
15: end while
16: Return  $\lfloor 2^{i-1} \rfloor$ 

```

---



# Theorem proof sketch

Therefore, the variance of  $S(h_{A,b}^i)$  is the sum of **individual** variances. And can be shown to be "small" compared to the mean.

By standard concentration inequalities,  $S(h_{A,b}^i)$  takes values close to it's mean reasonably often.

Computing median over  $T$  runs, guarantees an accurate estimate with high probability.

---

## Algorithm 9.1 ApproxModelCount ( $F, \mathcal{O}_S, \Delta$ )

---

```

1: Let  $S$  denote the set of solutions to the input formula  $F$ 
2:  $T \leftarrow \left\lceil \frac{\log(1/\Delta)}{\alpha} \log n \right\rceil$ 
3:  $i = 0$ 
4: while  $i \leq n$  do
5:   for  $t = 1, \dots, T$  do
6:     Sample hash function  $h_{A,b}^i : \Sigma \rightarrow \{0,1\}^i$ , i.e.
7:       sample uniformly  $A \in \{0,1\}^{i \times n}$ ,  $b \in \{0,1\}^i$ 
8:     Let  $S(h_{A,b}^i) = |\{x \in S \mid Ax \equiv b \pmod{2}\}|$ 
9:      $w_t^i \leftarrow \mathbb{I}[S(h_{A,b}^i) \geq 1]$ , using  $\mathcal{O}_S$  to check if  $\{x \in S \mid Ax \equiv b \pmod{2}\}$ 
10:   end for
11:   if Median( $w_1^i, \dots, w_T^i$ )  $< 1$  then
12:     break
13:   end if
14:    $i = i + 1$ 
15: end while
16: Return  $\lfloor 2^{i-1} \rfloor$ 

```

---



# Improving the approximation factor

Given the  $k$ -approximation algorithm and any  $\epsilon > 0$ , it's possible to design a  $(1 + \epsilon)$ -approximation algorithm, with the following construction:

Let  $l = \log_{1+\epsilon} k$ , define new set of configurations  $\sum^l = \sum \times \sum \dots \times \sum$ , and a new weight function  $w' : \sum^l \mapsto \mathcal{R}$  as  $w'(\sigma_1, \dots, \sigma_l) = w(\sigma_1)w(\sigma_2)\dots w(\sigma_l)$

Idea is to estimate the size of  $S(h_{A,b}^i)$  with multiple calls to an  $NP$ -oracle, e.g. check if  $S(h_{A,b}^i)$  contains at least  $k$  elements  $\rightarrow$  Reduces the variance, can be used to improve accuracy, but requires more calls to the  $NP$ -oracle

## Compactly representing constraints

- ▶ In model counting, the calls to an  $NP$ -oracle are implemented by using a SAT-solver. This is accomplished by adding to the formula  $i$  parity constraints, and checking for satisfiability.
- ▶ Naive encoding of a parity constraint over  $p$  variables requires  $2^{p-1}$  clauses of length  $p$  - ruling out  $2^p$  possible assignments (ones with wrong parity)
- ▶ Constraints can be compactly represented (introducing  $O(p)$  extra variables) using standard Tseitin transformation (Tseitin, 1983).

**Example:**  $x_1 \oplus x_2 \oplus x_3 \oplus x_4 = 0$  can be written as  
 $\{x_1 \oplus x_2 = z_1, x_2 \oplus x_3 = z_2, z_1 \oplus z_2 = 0\}$

# Practical implementations

- ▶ First practical implementation by Gomes et al. (2006) who used a SAT solver as an *NP*-oracle.
- ▶ Their algorithm leverages decades of research and engineering in SAT solving techniques for approximate model counting, resulted in huge improvements
- ▶ Recently, Chakraborty et al. (2013); Ivrii et al. (2015) provided several practical improvements
- ▶ Specifically, the former introduced the use of *pivots*, where an *NP*-oracle is used to check the existence of at least  $k > 1$  solutions ( $k = 1$  corresponds to earlier discussed algorithm), in order to improve the accuracy of the estimated count

# Probabilistic Models and Approx Inference: WISH Algorithm

- ▶ Generally, when the weight function is "close" for being constant on a subset of states, and zero elsewhere, then the hashing-based algorithm of Chakraborty et al. (2013) can be used - as in the model counting problem.
- ▶ Typical models in ML are unlikely to satisfy this restriction (e.g., weight function that is log-linear can have large variability)
- ▶ An alternative algorithm (WISH) based on universal hashing and combinatorial optimization which can handle general weight function, was introduced by Ermon et al. (2013)

# WISH Algorithm

---

**Algorithm 9.2** WISH ( $w : \Sigma \rightarrow \mathbb{R}^+, n = \log_2 |\Sigma|, \Delta, \alpha$ )

---

```
1:  $T \leftarrow \left\lceil \frac{\ln(n/\delta)}{\alpha} \right\rceil$ 
2: for  $i = 0, \dots, n$  do
3:   for  $t = 1, \dots, T$  do
4:     Sample hash function  $h_{A,b}^i : \Sigma \rightarrow \{0,1\}^i$ , i.e.
5:       sample uniformly  $A \in \{0,1\}^{i \times n}$ ,  $b \in \{0,1\}^i$ 
6:        $w_i^t \leftarrow \max_{\sigma} w(\sigma)$  subject to  $A\sigma \equiv b \pmod{2}$ 
7:   end for
8:    $M_i \leftarrow \text{Median}(w_i^1, \dots, w_i^T)$ 
9: end for
10: Return  $M_0 + \sum_{i=0}^{n-1} M_{i+1} 2^i$ 
```

---

Thank you for listening!

Any questions?